

Vandre.Tasks_I_II

November 20, 2017

```
In [1]: # Task I - Data Exploration
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import matplotlib.patches as mpatches
from sklearn import datasets, linear_model
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error

In [2]: plt.rcParams['figure.figsize'] = (12.0, 4.0)
X = pd.read_csv('../data/X.csv')
Y = pd.read_csv('../data/Y.csv')

In [3]: # 1. Compute the following statistics for each component of X and Y:
# mean, maximum, minimum, 25th percentile, 90th percentile, and standard
deviation.
metrics = pd.merge(X, Y, on='TimeStamp')
metrics.iloc[:,1:].describe(percentiles=[.25, .90])

Out[3]:
```

| | all_..idle | X..memused | proc.s | cswch.s | file.nr \ |
|-------|-------------|-------------|-------------|--------------|-------------|
| count | 3600.000000 | 3600.000000 | 3600.000000 | 3600.000000 | 3600.000000 |
| mean | 9.064981 | 89.137517 | 7.683303 | 54045.874022 | 2656.333333 |
| std | 16.122822 | 8.183662 | 8.532606 | 19497.811540 | 196.110748 |
| min | 0.000000 | 73.030000 | 0.000000 | 11398.000000 | 2304.000000 |
| 25% | 0.000000 | 82.965000 | 0.000000 | 31302.000000 | 2496.000000 |
| 50% | 0.000000 | 92.175000 | 6.000000 | 63908.000000 | 2688.000000 |
| 90% | 38.621000 | 96.770000 | 20.000000 | 72135.100000 | 2880.000000 |
| max | 69.540000 | 97.840000 | 48.000000 | 83880.000000 | 2976.000000 |

| | sum_intr.s | ldavg.1 | tcpsck | pgfree.s | DispFrames |
|-------|--------------|-------------|-------------|---------------|-------------|
| count | 3600.000000 | 3600.000000 | 3600.000000 | 3600.000000 | 3600.000000 |
| mean | 19978.040747 | 75.875772 | 48.997500 | 72872.154569 | 18.818394 |
| std | 4797.271325 | 43.862445 | 15.871155 | 19504.321175 | 5.219756 |
| min | 10393.000000 | 11.130000 | 21.000000 | 15928.000000 | 0.000000 |
| 25% | 16678.000000 | 28.200000 | 34.000000 | 61601.750000 | 13.390000 |
| 50% | 18109.000000 | 75.390000 | 47.000000 | 71686.500000 | 19.120001 |
| 90% | 28228.400000 | 127.993000 | 71.000000 | 97532.500000 | 24.610000 |
| max | 35536.000000 | 147.470000 | 87.000000 | 145874.000000 | 30.390000 |

In [4]: # 2. Compute the following quantities:

```
# (a) the number of observations with memory usage larger than 80%;
a = metrics[metrics['X..memused'] > 80].count()['X..memused']
print 'a) Number of observations:', a


# (b) the average number of used TCP sockets for observations
# with more than 18000 interrupts/sec;
b = metrics[metrics['sum_intr.s'] > 18000].mean()['tcpsck']
print 'b) Average number of used TCP sockets: %.2f' % b

# (c) the minimum memory utilization for observations
# with CPU idle time lower than 20%.
c = metrics[metrics['all_..idle'] < 20].min()['X..memused']
print 'c) Minimum memory utilization: %.2f' % c
```

a) Number of observations: 2875
b) Average number of used TCP sockets: 46.35
c) Minimum memory utilization: 73.03


In [5]: # 3. Produce the following plots:

```
# (a) Time series of percentage of idle CPU and of used memory
# (both in a single plot);
series = pd.DataFrame(metrics, columns=['all_..idle', 'X..memused']);
ts = series.plot(title = 'Time series of idle CPU and used memory')
ts.legend(['Idle CPU', 'Used memory'])
ts.set_xlabel("Second")
ts.set_ylabel("Percentage (%)")
fig = ts.get_figure()
plt.show()
```




```
In [6]: plt.close(fig)
```

```
In [7]: # (b) Density plots, histograms, and box plots of idle CPU and of used memory.  
kde = series.plot.kde(title = 'Density of idle CPU and used memory')  
kde.legend(['Idle CPU', 'Used memory'])  
kde.set_xlabel("Percentage (%)")  
fig = kde.get_figure()  
plt.show()
```



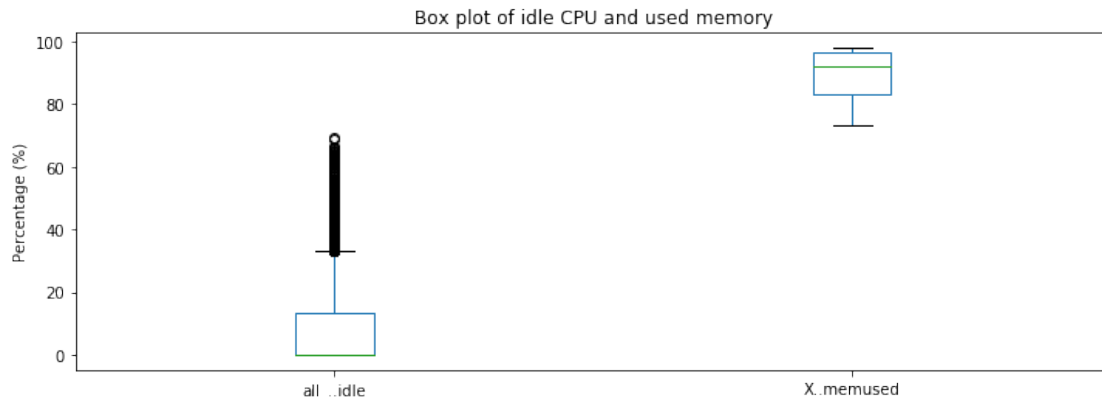
```
In [8]: plt.close(fig)
```

```
In [9]: hist = series.plot.hist(title = 'Histogram of idle CPU and used memory')  
hist.legend(['Idle CPU', 'Used memory'])  
hist.set_xlabel("Percentage (%)")  
fig = hist.get_figure()  
plt.show()
```



```
In [10]: plt.close(fig)
```

```
In [11]: bp = series.plot.box(title = 'Box plot of idle CPU and used memory')
bp.set_ylabel("Percentage (%)")
fig = bp.get_figure()
plt.show()
```



```
In [12]: plt.close(fig)
```

```
In [13]: # Task II - Estimating Service Metrics from Device Statistics
```

```
def nmae(y_test, y_pred):
    y_test = pd.DataFrame(y_test)
    y_pred = pd.DataFrame(y_pred)

    sum = 0.0
    m = len(y_test)
    average = np.mean(y_test)

    for i in range(m):
        sum += abs(y_test.iloc[i, 0] - y_pred.iloc[i, 0])
        i += 1

    nmae = (sum / m) / average
    return nmae
```

```
In [14]: # 1. Evaluate the Accuracy of Service Metric Estimation
```

```
# (a) Model Training - use linear regression to train a model M
# with the training set. Provide the coefficients of your model M.
X = X.iloc[:, X.columns != 'TimeStamp']
Y = Y.iloc[:, Y.columns != 'TimeStamp']

X_train, X_test, y_train, y_test = train_test_split(X, Y, train_size=0.7,
test_size=0.3)
```

```

lm = linear_model.LinearRegression()
model = lm.fit(X_train, y_train)
y_pred = lm.predict(X_test)

```

```

print 'Coefficients:', lm.coef_

```

```

Coefficients: [[ -8.63879476e-02  -7.72699183e-02  -1.33623558e-02
 -9.87963259e-05
 -3.54053585e-03   2.15181648e-05  -5.89256394e-02  -5.60887870e-02
 -1.58745163e-05]]

```

```

In [15]: # (b) Accuracy of Model M - compute the estimation error of M over the test set.
# Explained variance score: 1 is perfect prediction
print 'Accuracy score: %.4f' % model.score(X_test, y_test)
print 'Normalized Mean Absolute Error: %.4f ' % nmae(y_test, y_pred)

```

```

Accuracy score: 0.7282

```

```

Normalized Mean Absolute Error: 0.1012

```

```

In [16]: plt.clf()

```

```

In [17]: # (c) Produce a time series plot that shows both the measurements and
# the model estimations for M for the Video Frame Rate values in the
# test set (see example of such a plot in Figure 4(a) of [1]).
# Show also the prediction of the a naive method.
interval = np.arange(0, 1080, 1)

```

```

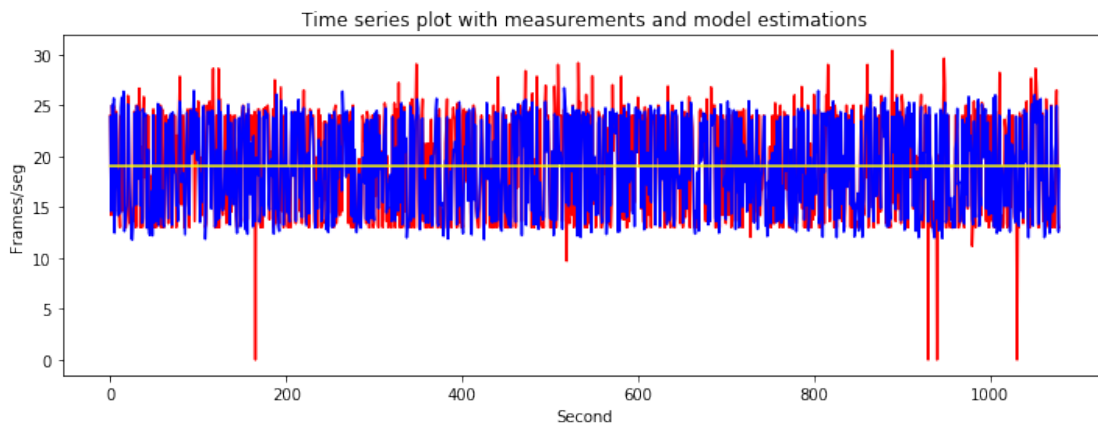
naive_pred = [np.mean(y_test) for i in range(len(y_test))]
naive_pred = pd.DataFrame(naive_pred)

```

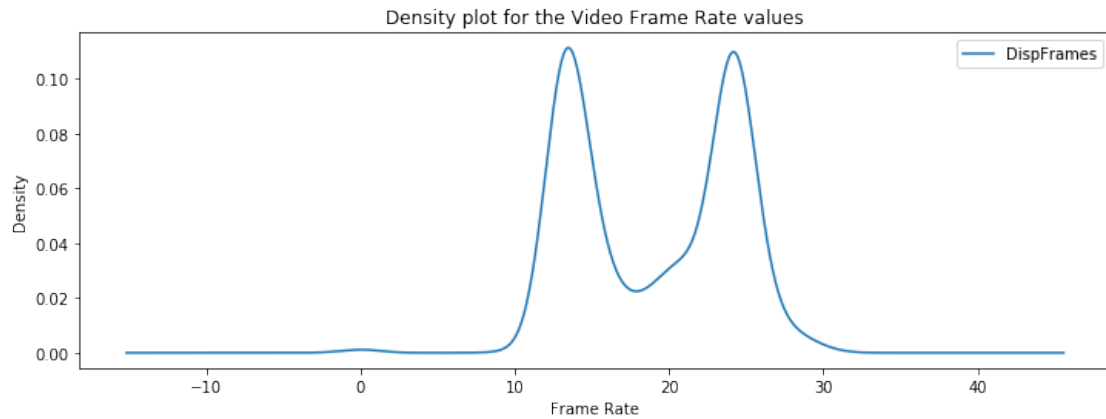
```

plt.plot(interval, y_test, 'red', interval, y_pred, 'blue', interval,
naive_pred, 'yellow')
plt.title('Time series plot with measurements and model estimations')
plt.ylabel('Frames/seg')
plt.xlabel('Second')
plt.show()

```




```
In [18]: # (d) Produce a density plot and a histogram for the Video Frame Rate values
# in the test set. Set the bin size of the histogram to 1 frame.
kde = y_test.plot.kde(title = 'Density plot for the Video Frame Rate values')
kde.set_xlabel("Frame Rate")
fig = kde.get_figure()
plt.show()
```



```
In [19]: plt.close(fig)
```

```
In [20]: plt.hist(y_test.DispFrames, bins=1, normed=1)
plt.axis([0, 60, 0, 0.1])
fig = hist.get_figure()
plt.show()
```




```
In [21]: plt.close(fig)
```

```
In [22]: # (e) Produce a density plot for the prediction errors in the test set.
dy = pd.DataFrame(y_test)
```

```
residuals = [dy.iloc[i] - y_pred[i] for i in range(len(y_pred))]
residuals = pd.DataFrame(residuals)
```

```
kde = residuals.plot.kde(title = 'Density plot for prediction errors')
fig = kde.get_figure()
plt.show()
```



```
In [23]: # (f) Based on the above figures and graphs, discuss the accuracy of
# estimating the Video Frame Rate.
```

```
# Considering the figures and graphs above, especially the time series plot,
# it's noticeable that estimates are not so accurate but generally close to
# the real frame rate. On the other hand, the deviation is usually high when
# a frame rate drop happens.
```

```
In [24]: plt.close(fig)
plt.clf()
```

```
In [25]: # 2. Study the Relationship between Estimation Accuracy and the Size of
# the Training Set (a) From the above training set with 2520 observations,
# create six training sets by selecting uniformly at random 50, 100, 200,
# 500, 1000, and 2520 observations (which is the original set).
# (b) Train a linear model and compute the NMAE for each model for the
# original test set with 1080 observations. (c) Perform the above 50 times,
# so you train models for 50 different subsets of a given size.
```

```
NMAE = {}
group_sizes = [2520, 1000, 500, 200, 100, 50]
rounds = 50
```

```

for size in group_sizes:
    group_results = []

    for i in range(rounds):
        if(size == 2520):
            gX_train = X_train
            gy_train = y_train
        else:
            gX_train, gX_test, gy_train, gy_test = train_test_split(X_train,
y_train, train_size=size, test_size=0)

        lm = linear_model.LinearRegression()
        model = lm.fit(gX_train, gy_train)
        y_pred = lm.predict(X_test)

        res = nmae(y_test, y_pred)
        group_results.append([i, res])

    NMAE[size] = np.array(group_results)

```

In [26]: *# (d) Produce a plot that shows NMAE for M against the size of
the training set. Use error bars or box plots to show the
range of the NMAE values for a given set size.*

```

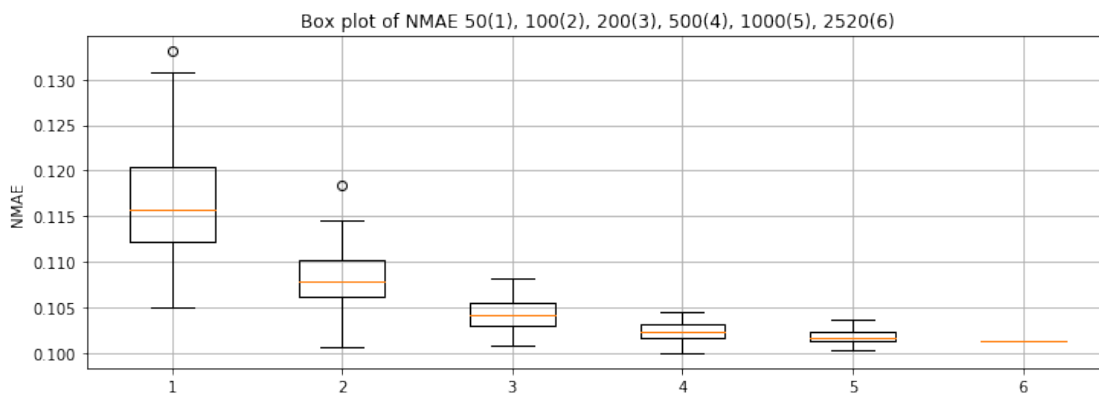
data = [NMAE[50][:,1], NMAE[100][:,1], NMAE[200][:,1], NMAE[500][:,1],
NMAE[1000][:,1], NMAE[2520][:,1]]

```

```

plt.title('Box plot of NMAE 50(1), 100(2), 200(3), 500(4), 1000(5), 2520(6)')
plt.ylabel('NMAE');
plt.grid(True)
plt.boxplot(data)
plt.show()

```



In [27]: # (e) Based on the above, discuss the relationship between the accuracy
of the model estimations and the training set.

Estimations are more accurate when the amount of observations in
a training set is closer to the total amount of observations.